

DSI-06_Homework 5: Chapter 6, pg 286

Julia Gallucci

2023-03-05

13. In this exercise, we will predict the number of applications received using the other variables in the College data set.

```
#install.packages("ISLR") install package containing College dataset  
library(ISLR) #load library  
attach(College) #attach College dataset to make the variables associated with College available.  
head(College) #return the column names and first few rows of the dataset
```

##	Private	Apps	Accept	Enroll	Top10perc	Top25perc
## Abilene Christian University	Yes	1660	1232	721	23	52
## Adelphi University	Yes	2186	1924	512	16	29
## Adrian College	Yes	1428	1097	336	22	50
## Agnes Scott College	Yes	417	349	137	60	89
## Alaska Pacific University	Yes	193	146	55	16	44
## Albertson College	Yes	587	479	158	38	62

##	F.Undergrad	P.Undergrad	Outstate	Room.Board	Books
## Abilene Christian University	2885	537	7440	3300	450
## Adelphi University	2683	1227	12280	6450	750
## Adrian College	1036	99	11250	3750	400
## Agnes Scott College	510	63	12960	5450	450
## Alaska Pacific University	249	869	7560	4120	800
## Albertson College	678	41	13500	3335	500

##	Personal	PhD	Terminal	S.F.Ratio	perc.alumni	Expend
## Abilene Christian University	2200	70	78	18.1	12	7041
## Adelphi University	1500	29	30	12.2	16	10527
## Adrian College	1165	53	66	12.9	30	8735
## Agnes Scott College	875	92	97	7.7	37	19016
## Alaska Pacific University	1500	76	72	11.9	2	10922
## Albertson College	675	67	73	9.4	11	9727

##	Grad.Rate
## Abilene Christian University	60
## Adelphi University	56
## Adrian College	54
## Agnes Scott College	59
## Alaska Pacific University	15
## Albertson College	55

(a) Split the data set into a training set and a test set.

We can use the function `sample()` to split the observations into two halves for the training and testing sets. The train vector contains the indices of the College dataset that will be used as the training data.

```
# split the 777 observations in half (roughly) so 388 observations in the training set
train <- sample(1:nrow(College), nrow(College) / 2)
test <- (-train)
```

(b) Fit a linear model using least squares on the training set, and report the test error obtained.

```
lm.fit <- lm(Apps ~ ., data = College, subset = train) # fit model on only training data
lm.pred <- predict(lm.fit, College[test,], type="response") #predict using our fit model on test data
mean((lm.pred - College[test,]$Apps)^2) #given our predictions we can compute the MSE for the test set
```

```
## [1] 1135758
```

So our estimated test MSE for the linear regression is 1135758

(c) Fit a ridge regression model on the training set, with lambda chosen by cross-validation. Report the test error obtained.

We will use the glmnet() function from the glmnet package to perform ridge regression. Instead of using a data frame, the glmnet() function requires that the predictors be in the form of a matrix and the response be in the form of a vector. The model.matrix() automatically transforms qualitative variables into dummy variables.

```
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-6
```

```
#Set up matrices needed for the glmnet functions
train.mat = model.matrix(Apps~., data = College[train,]) #training matrix using train observations
test.mat = model.matrix(Apps~., data = College[test,]) #testing matrix using test observations
```

Choose lambda using cross-validation

```
lamda_cv = cv.glmnet(train.mat, College[train,]$Apps, alpha=0) #perform cross validation on our model
bestlam = lamda_cv$lambda.min #minimum lambda extracted
bestlam
```

```
## [1] 405.8404
```

Fit a ridge regression model and predict on our testing data using our best lambda!

```
ridge.mod = glmnet(train.mat, College[train,]$Apps, alpha = 0)
#Make predictions
ridge.pred = predict(ridge.mod, s=bestlam, newx = test.mat)

mean((ridge.pred - College[test,]$Apps)^2) #given our predictions we can compute the MSE for the test set

## [1] 976261.5
```

The test error of the ridge regression fit with a lambda chosen by cross-validation is 976261.5, lower than the linear model test error!

(d) Fit a lasso model on the training set, with lambda chosen by crossvalidation. Report the test error obtained, along with the number of non-zero coefficient estimates.

We will perform the lasso the exact same way as ridge regression except we have $\alpha = 1$ in the `glmnet()` function to specify the lasso method is to be used.

Choose lambda using cross-validation

```
lamda_cv2 = cv.glmnet(train.mat,College[train,]$Apps,alpha=1)
bestlam_2 = lamda_cv2$lambda.min
bestlam_2
```

```
## [1] 1.97344
```

Fit a lasso regression model and predict on our testing data using our best lambda!

```
#Fit lasso model
lasso.mod = glmnet(train.mat,College[train,]$Apps,alpha = 1)
#Make predictions
lasso.pred = predict(lasso.mod,s=bestlam_2,newx=test.mat)
mean((lasso.pred - College[test,]$Apps)^2)
```

```
## [1] 1115901
```

The test error of the lasso model fit with a lambda chosen by cross-validation is 1115901 ,higher thanridge regression test error but lower than the linear model!

Based on test error, the model performance from best to worst is:

1. Ridge Regression (976261.5)
2. Lasso model (1115901)
3. Linear model (1135758)